

Jaeyong Song

Accelerated Intelligent Systems Lab. (AISys) @ Seoul National University

Email: jaeyong.song@snu.ac.kr

Profile: <https://aisys.snu.ac.kr/members/JaeyongSong.html>

Google scholar: https://scholar.google.com/citations?user=2H-m_VYAAAAJ

LinkedIn: <https://www.linkedin.com/in/jyong-song/>

Research Interest: SYSTEMS & ARCHITECTURES FOR DEEP LEARNING

Distributed Deep Learning, Accelerators, Large Language Models (LLMs), Graph Neural Networks (GNNs)

EDUCATION

- | | |
|---|---------------------------------------|
| Seoul National University | Seoul, Korea |
| • <i>Ph.D. Candidate - Electrical & Computer Engineering; Current GPA: 4.20/4.3</i> | <i>September 2023 - February 2027</i> |
| Yonsei University | Seoul, Korea |
| • <i>M.S. - Computer Science; GPA: 4.25/4.3</i> | <i>September 2021 - August 2023</i> |
| Yonsei University | Seoul, Korea |
| • <i>B.E. - Applied Statistics, B.S. - Computer Science; GPA: 4.17/4.3 (3.96/4.0, Rank: 2/88)</i> | <i>- August 2021</i> |

PUBLICATIONS

¹: indicates co-first authors.

- **GriNNder: Breaking the Memory Capacity Wall in Full-Graph GNN Training with Storage Offloading:** Jaeyong Song, Seongyeon Park, Hongsun Jang, Jaewon Jung, Hunseong Lim, Junguk Hong, Jinho Lee *MLSys*, 2026
- **A Cost-Effective Near-Storage Processing Solution for Offline Inference of Long-Context LLMs:** Hongsun Jang, Jaeyong Song, Changmin Shin, Si Ung Noh, Jaewon Jung, Jisung Park, Jinho Lee *ASPLOS*, 2026
- **FlexiWalker: Extensible GPU Framework for Efficient Dynamic Random Walks with Runtime Adaptation:** Seongyeon Park, Jaeyong Song, Changmin Shin, Sukjin Kim, Junguk Hong, Jinho Lee *EuroSys*, 2026
- **FALA: Locality-Aware PIM-Host Cooperation for Graph Processing with Fine-Grained Column Access:** Changmin Shin, Jaeyong Song, Seongmin Na, Jun Sung, Hongsun Jang, Jinho Lee *MICRO*, 2025
- **Piccolo: Large-Scale Graph Processing with Fine-Grained In-Memory Scatter-Gather:** Changmin Shin, Jaeyong Song, Hongsun Jang, Dogeun Kim, Jun Sung, Taehee Kwon, Jae Hyung Ju, Frank Liu, Yeonkyu Choi, and Jinho Lee *HPCA*, 2025
- **GraNNDis: Fast Distributed Graph Neural Network Training Framework for Multi-Server Clusters:** Jaeyong Song, Hongsun Jang, Hunseong Lim, Jaewon Jung, Youngsok Kim, and Jinho Lee *PACT*, 2024
- **A Case for In-Memory Random Scatter-Gather for Fast Graph Processig:** Changmin Shin¹, Taehee Kwon¹, Jaeyong Song, Jae Hyung Ju, Frank Liu, and Jinho Lee *IEEE CAL*, 2024
- **PeerAiD: Improving Adversarial Distillation from a Specialized Peer Tutor:** Jaewon Jung, Hongsun Jang, Jaeyong Song, and Jinho Lee *CVPR*, 2024
- **AGATHA: Fast and Efficient GPU Acceleration of Guided Sequence Alignment for Long Read Mapping:** Seongyeon Park, Junguk Hong, Jaeyong Song, Hajin Kim, Youngsok Kim, and Jinho Lee *PPoPP*, 2024
- **Pipette: Automatic Fine-grained Large Language Model Training Configurator for Real-World Clusters:** Jinkyu Yim¹, Jaeyong Song¹, Yerim Choi, Jaeben Lee, Jaewon Jung, Hongsun Jang, and Jinho Lee *DATE*, 2024
- **Smart-Infinity: Fast Large Language Model Training using Near-Storage Processing on a Real System:** Hongsun Jang, Jaeyong Song, Jaewon Jung, Jaeyoung Park, Youngsok Kim, and Jinho Lee *HPCA (Best Paper Award Honorable Mention)*, 2024
- **Fast Adversarial Training with Dynamic Batch-level Attack Control:** Jaewon Jung, Jaeyong Song, Hongsun Jang, Hyeyoon Lee, Kanghyun Choi, Noseong Park, and Jinho Lee *DAC*, 2023
- **Pipe-BD: Pipelined Parallel Blockwise Distillation:** Hongsun Jang, Jaewon Jung, Jaeyong Song, Joonsang Yu, Youngsok Kim, and Jinho Lee *DATE*, 2023
- **Optimus-CC: Efficient Large NLP Model Training with 3D Parallelism Aware Communication Compression:** Jaeyong Song¹, Jinkyu Yim¹, Jaewon Jung, Hongsun Jang, Hyung-Jin Kim, Youngsok Kim, and Jinho Lee *ASPLOS*, 2023
- **SGCN: Exploiting Compressed-Sparse Features in Deep Graph Convolutional Network Accelerators:** Mingi Yoo¹, Jaeyong Song¹, Jounghoo Lee, Namhyung Kim, Youngsok Kim, and Jinho Lee *HPCA*, 2023
- **Slice-and-Forge: Making Better Use of Caches for Graph Convolutional Network Accelerator:** Mingi Yoo¹, Jaeyong Song¹, Hyeyoon Lee, Jounghoo Lee, Namhyung Kim, Youngsok Kim, and Jinho Lee *PACT (Best Paper)*, 2022
- **Making a Better Use of Caches for GCN Accelerators with Feature Slicing and Automatic Tile Morphing:** Mingi Yoo¹, Jaeyong Song¹, Jounghoo Lee, Namhyung Kim, Youngsok Kim, and Jinho Lee *IEEE CAL*, 2021

HONORS AND AWARDS

- **Best Paper Award** @ PACT '22 (Slice-and-Forge) - Oct. 2022
- **Best Paper Award Honorable Mention** @ HPCA '24 (Smart-Infinity) - Mar. 2024
- Encouragement Prize (Top 6%) @ The 30th Samsung Humantech Paper Award - Feb. 2024
- Excellence Prize @ LG Display AI & Big Data Competition - Dec. 2022
- Yonsei Social Entrepreneurship Award - Database-based learning management system development - Jan. 2020
- Academic (Highest) Honors (Undergraduate, four times) - 2016-1 (honors), 2019-1 (honors), 2019-2 (highest honors), 2021-1 (honors) - GPA: 3.96/4.0 (4.17/4.3, Rank: 2/88)
- Commendation from Jeju Provincial Police Agency - Commendation from the local police commissioner - Oct. 2017

TEACHING EXPERIENCE

• Embedded Systems Design	Seoul National University
• <i>Teaching Assistant</i>	<i>Fall 2024</i>
• Digital Systems Design and Lab.	Seoul National University
• <i>Teaching Assistant</i>	<i>Fall 2023</i>
• Logic Circuit Design	Yonsei University
• <i>Teaching Assistant</i>	<i>Fall 2021</i>